RAG ASSIGNMENT REPORT / 19 OCTOBER 2026

Baseline Retrieval and Reranking



Dense retrieval and Cross-Encoder evidence selection

Varun Patil 2025204013 Code and saved results

ABSTRACT

Cross-Encoder reranking restored decisive evidance for three selected questions that baseline retrieval could not answer from its first three chunks. The experiment uses NIST SP 800-63B-4, *Digital Identity Guidelines: Authentication and Authenticator Management*, final July 2025 [1]. Two controls shows why an improved rank need not improve an answer, and why missing candidates limit reranking.

1. Architecture comparison

The 129-page PDF yields **409 chunks** from PDF pages 12–129, excluding front matter. Each page is divided into 180-token windows with 40-token overlap. Both pipelines uses this same corpus. The Bi-Encoder, `all-MiniLM-L6-v2`, embeds chunks and queries independantly [2]. Chunk embeddings are precomputed; both query and chunk vectors are normalized. Exact FAISS inner-product search is therefore equivalent to cosine similarity.

Two retrieval paths

Baseline: query embedding → Top 10 search → original Top 3 → Gemini.

Advanced: same Top 10 → Cross-Encoder → new Top 3 → Gemini.

Baseline supplies the original first three candidates to Gemini. The Cross-Encoder, `ms-marco-MiniLM-L6-v2`, jointly processes all ten query–chunk pairs, allowing self-attention across their tokens [3, 4]. It sorts by descending relevance score and selects a new Top 3. Its scores are neither probabilities nor cosine similarities.

Both answers use `gemini-3.8-flash`, temperature 0, seed 42, and the same evidence-and-citation prompt template [5]. Only the selected excerpts change. These settings do not guarantee deterministic generation.

2. Qualitative and rank shift analysis

Five cases was selected from 44 exploratory queries. The three recoveries were checked for equivalent evidence across the full baseline context. They are selected stress tests, not an estimate of average accuracy. The six answers below are reproduced in three pairs.

D10 Resetting failed attempts

Rank 4 → 1

Failure type: Negation and constraints

Question. When a user requests a fresh out-of-band verification code, can the failed-attempt counter start over?

PDF 33 / printed 22; section 3.1.3.2. Tracked chunk: `p033_c03`.

Source excerpt [1]. *"Generating a new authentication secret SHALL NOT reset the failed authentication count."*

Answer 1: Baseline

The provided excerpts do not establish whether a failed-attempt counter can start over when a user requests a fresh out-of-band verification code. The excerpts discuss out-of-band authenticators and recovery code requirements, including validity periods, format, and throttling under Sec. 3.2.2 [p059_c03, p031_c04, p059_c02], but they completely lack any mention of a failed-attempt counter or the conditions under which it may be reset.

Answer 2: Reranked

No. Generating a new authentication secret SHALL NOT reset the failed authentication count [p033_c02, p033_c03].

Interpretation. Baseline returned recovery-code material rather than the counter restricton. Reranking changes an explicit abstention into a supported answer; it does not correct a fabricated reset rule.

D28 Reusing biometric samples

Rank 5 → 2

Failure type: Document-specific permission

Question. Can face images collected during sign-in be reused to improve the matching template or for research, and what permission is needed?

PDF 45 / printed 34; section 3.2.3.4. Tracked chunk: `p045_c01`.

Source excerpt [1]. *"Biometric samples collected in the authentication process MAY be used locally within the authenticator or verifier performing the biometric comparison to update the templates for the purpose of compensating for changes in subscriber characteristics or — with explicit subscriber consent — other research purposes."*

Answer 1: Baseline

The provided excerpts do not establish whether face images collected during sign-in can be reused to improve the matching template or for research, nor do they establish what permission is needed.

The excerpts only state that facial images do not constitute secrets and can often be obtained without consent [p043_c01], detail usability considerations that affect facial recognition accuracy (e.g., artifacts, lighting, expressions, poses) [p089_c01], and define central comparison requirements—including that sensors/endpoints SHALL be authenticated, appropriate controls SHALL be implemented for sensitive personal information, and verifiers MAY limit centrally stored templates to particular sensors [p044_c04]. Any provisions regarding the reuse of sign-in biometric samples or required permissions are missing.

Answer 2: Reranked

Biometric samples collected during the authentication process MAY be used locally within the authenticator or verifier performing the biometric comparison to update templates (to compensate for changes in subscriber characteristics) or for other research purposes [p045_c01].

Regarding permission:

- For research purposes, "explicit subscriber consent" is required [p045_c01].
- For updating templates to compensate for changes in subscriber characteristics, the text does not state that permission is needed [p045_c01].

Additionally, the samples and any derived biometric data SHALL be erased immediately after adaptation or research data has been derived, and an allowable time limit for adaptation SHALL be established and enforced by the authenticator or CSP [p045_c01].

Interpretation. Baseline lack the reuse and consent clause. Reranking restores local use and explicit subscriber consent for other reserach. The provision does not state that template adaptation needs consent; this does not establish an exemption from all other applicable consent requirements.

D34 Sending a postal recovery notice

Rank 6 → 1

Failure type: Exception and corner case

Question. If the only electronic notification address was also used to receive a recovery code, can the provider skip sending a postal recovery notice?

PDF 63 / printed 52; section 4.6. Tracked chunk: `p063_c02`.

Source excerpt [1]. *"Account recovery notifications SHALL also be sent to a postal address if the only other notification address in the subscriber account is the address to which an issued recovery code was sent."*

Answer 1: Baseline

The provided excerpts do not establish the answer. They do not contain any information or requirements regarding postal recovery notices, nor do they state whether a provider can skip sending a postal recovery notice if the only electronic notification address was used to receive a recovery code [p059_c02, p059_c03, p060_c01].

Answer 2: Reranked

No, the provider cannot skip sending a postal recovery notice [p063_c02].

According to the excerpt, account recovery notifications SHALL also be sent to a postal address if the only other notification address in the subscriber account is the address to which an issued recovery code was sent [p063_c02].

Interpretation. Baseline cover recovery-code delivery, omitting the notification exception. Reranking restores the seperate postal-notice obligation when the only other notification address received the issued recovery code.

Controls and retrieval limits

D03 / equivalent evidence. The AAL3 prohibition for cloud-synced passkeys, `p020_c01` (PDF 20, printed 9, section 2.3.2), moves from rank 4 to 2. But baseline rank 2 already contains equivalent evidence in `p103_c02` (PDF 103, printed 92). Both answers correctly reject synced passkeys at AAL3. This is not a failure recovery.

D05 / candidate recall. The voice-biometric prohibition in `p043_c04` (PDF 43, printed 32, section 3.2.3.2) is absent from the initial Top 10. Baseline incorrectly permits voice biometrics under general biometric conditions; the reranked answer abstains. This avoids unsupported permission without recovering the missing rule. Full control answers and candidate lists remain in the notebook and results.

3. Latency profiling

Each query has one warm-up and seven measured repetitons: **35 samples**. The run uses Python 3.12.3 on an ARM64 Mac with macOS 26.6.2, CPU inference, and one thread each for Torch and FAISS. Table 1 reports medians from `results/latency_summary.csv`.

TABLE 1 / RETRIEVAL LATENCY

Case	Baseline (ms)	Reranking increment (ms)	Advanced total (ms)	Total / baseline
D10	4.70	127.72	132.14	28.11×
D28	4.76	125.96	130.72	27.46×
D34	4.64	125.72	130.52	28.16×
D03	4.68	124.72	129.50	27.69×
D05	4.85	133.90	138.65	28.60×

Medians in milliseconds. Model loading, PDF processing, indexing, and Gemini generation are excluded.

Baseline includes query embedding, vector search, and candidate construction. The reranking increment includes pair validation, tokenization, inference, sorting, and result construction. Advanced total adds both stages *within each repetition*; separate medians need not sum exactly.

Advanced retrieval takes about **27.5–28.6 times as long** here. Seven samples offers limited precision. The latest run recomputed retrieval and timings but reused all ten earlier Gemini responses from an exact prompt/model/configuration cache, with **zero fresh API answers**. Original response timestamps are retained in `results/answers.json`.

4. Why reranking follows retrieval

A Cross-Encoder must jointly process every query–chunk pair. Scoring the whole corpus would require **409 pairs rather than ten**: 40.9 times as many pairs, *not* a measured 40.9-fold latency increase. Bi-Encoder document vectors are reused across queries; the vector store restricts expensive joint scorring to ten candidates. This implementation uses exact FAISS search, rather than approximate search.

Pair counts grows with the collection. D05 exposes the trade-off: an initial retrieval miss prevents reranking from restoring the absent rule. Reranking improves selection only within the candidates it receives.

Overall takeaway

Across the three selected recovery cases, reranking restored the missing rule and changed an abstention into a supported answer. These examples shows the benefit of better evidence selection, while the median reranking increment across all five cases was about 125–134 ms per query. The D05 control also shows that reranking cannot repair an initial Top 10 miss, so first-stage retrival quality still matters.

References

[1] NIST. SP 800-63B-4: Authentication and Authenticator Management. Final, July 2025.

[2] Sentence Transformers. `sentence-transformers/all-MiniLM-L6-v2` model card.

[3] Cross-Encoder. `cross-encoder/ms-marco-MiniLM-L6-v2` model card.

[4] Sentence Transformers. Cross-Encoder usage documentation.

[5] Google AI for Developers. Gemini content generation documentation.